

PARALLEL ALGORITHM FOR UPDATING SINGLE-SOURCE SHORTEST PATHS IN DYNAMIC NETWORKS

Presented by:

- Haris Ahmed
- Aazib Abdullah
- Amir akbar

INTRODUCTION TO PROBLEM DOMAIN

- Graphs are essential in modelling real-world networks
- Single-Source Shortest Path (SSSP) problems are core in graph analysis.
- Real-world networks (social, traffic) constantly change
- Traditional recomputation is costly and inefficient

PROBLEM STATEMENT

- Dynamic updates to graphs disrupt shortest paths.
- Traditional algorithms like Dijkstra's don't scale well on updates.
- Need for a parallel, platform-independent update framework
- Avoid recomputation by updating affected subgraphs

SSSP PROBLEM & CHALLENGES

- Goal: Find shortest paths from one source to all nodes
- Dynamic graphs: edge insertions and deletions
- Key issues: synchronization, load imbalance, performance

FRAMEWORK OVERVIEW

- Two-step update strategy:
 - Identify affected subgraph
 - Update SSSP tree locally
- Parallel and asynchronous processing
- Works without locks or critical sections

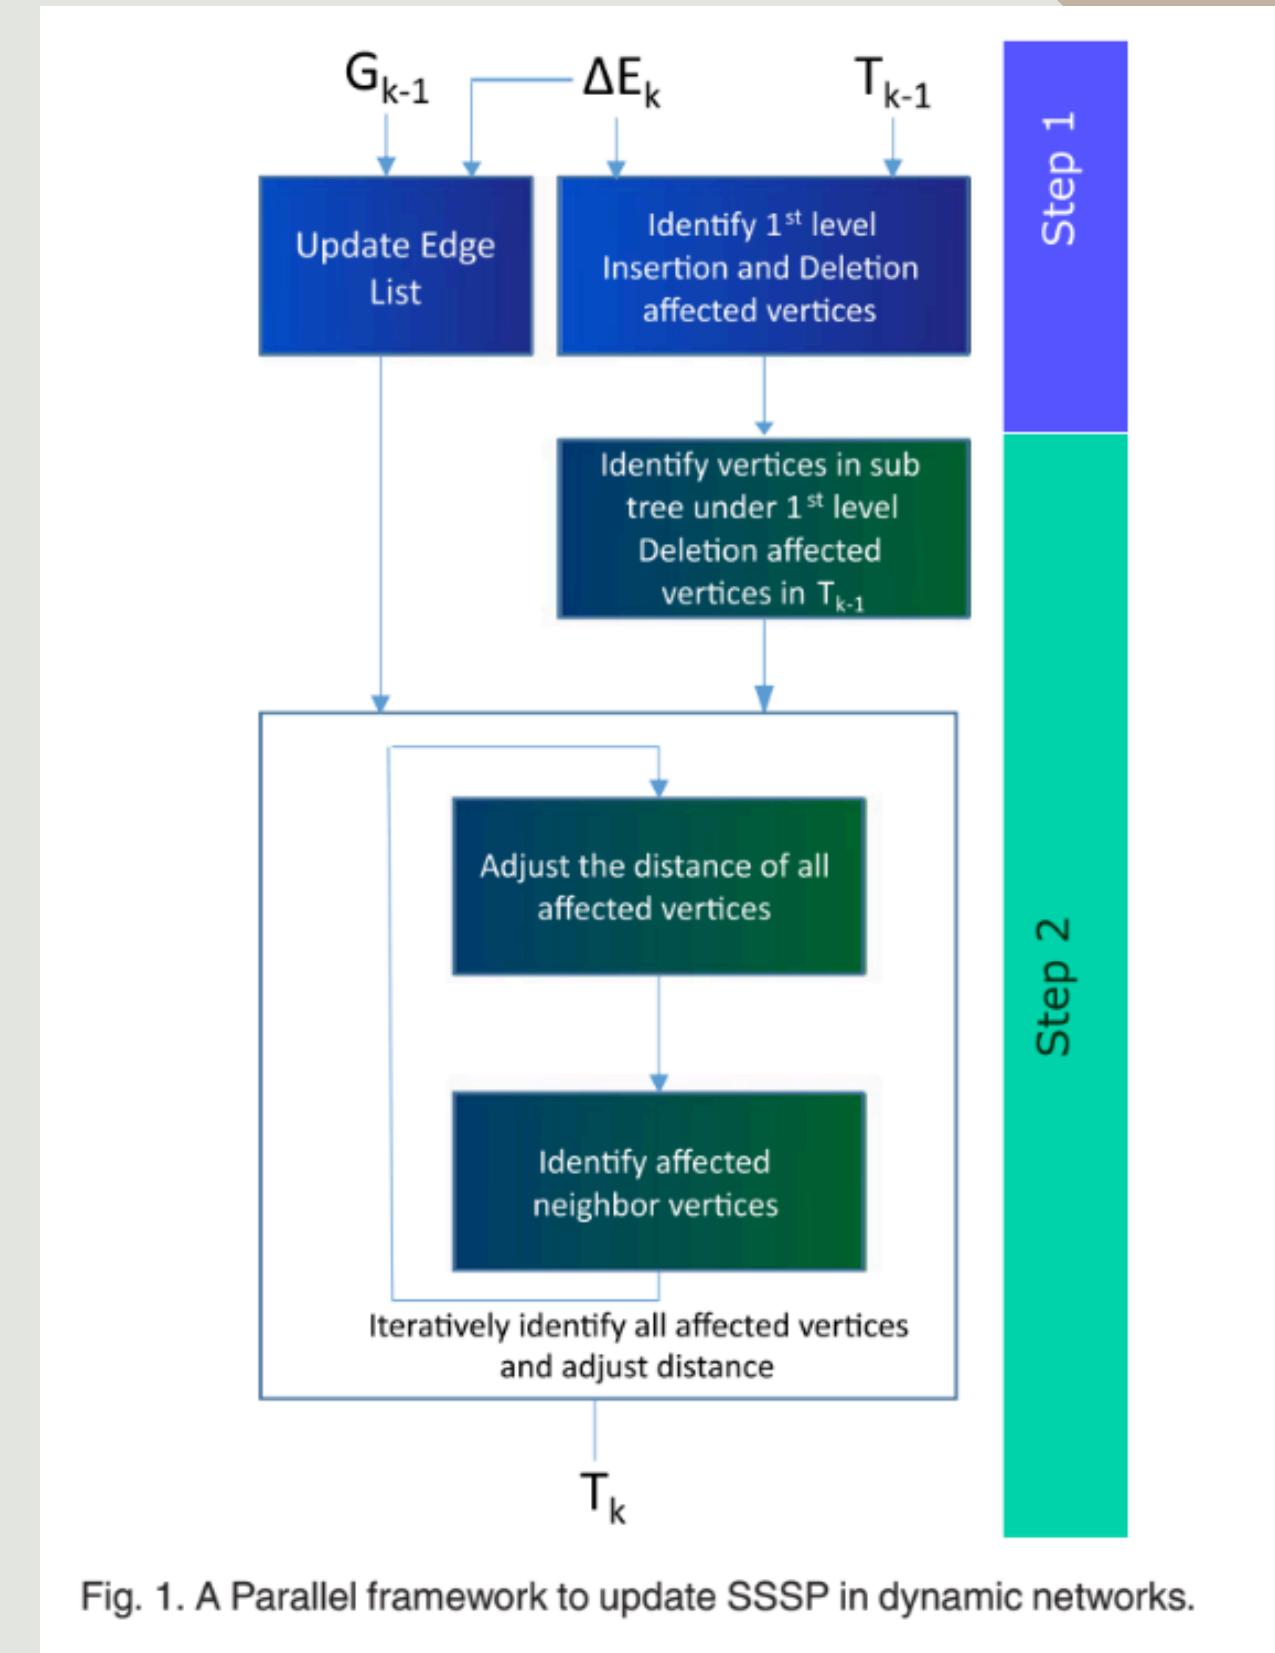


Fig. 1. A Parallel framework to update SSSP in dynamic networks.

STEP 1 – IDENTIFYING AFFECTED SUBGRAPH

- Parallel processing of changed edges
- Deletions: disconnect affected subtrees
- Insertions: update if new shorter path is found
- Affected vertices marked using flags

Algorithm 2. Identify Affected Vertices

```
1: Function ProcessCE ( $G, T, \Delta E, Parent, Dist$ ):  
2:   Initialize Boolean arrays  $Affected\_Del$  and  $Affected$  to  $false$   
3:    $G_u \leftarrow G$   
4:   for each edge  $e(u, v) \in Del_k$  in parallel do  
5:     if  $e(u, v) \in T$  then  
6:        $y \leftarrow argmax_{x \in \{u, v\}}(Dist[x])$   
7:       Change  $Dist[y]$  to infinity  
8:        $Affected\_Del[y] \leftarrow True$   
       $Affected[y] \leftarrow True$   
9:       Mark  $e(u, v)$  as deleted  
10:      for each edge  $e(u, v) \in Ins_k$  in parallel do  
11:        if  $Dist[u] > Dist[v]$  then  
12:           $x \leftarrow v, y \leftarrow u$   
13:        else  
14:           $x \leftarrow u, y \leftarrow v$   
15:        if  $Dist[y] > Dist[x] + W(x, y)$  then  
16:           $Dist[y] \leftarrow Dist[x] + W(x, y)$   
17:           $Parent[y] \leftarrow x$   
18:           $Affected[y] \leftarrow True$   
19:        Add  $e(u, v)$  to  $G_u$ 
```

STEP 2 – UPDATING SUBGRAPH

- Update affected vertices iteratively
- Traverse neighbors to update distances and parents
- Uses convergence logic instead of synchronization
- Ends when no updates remain

Algorithm 3. Update Affected Vertices

```
1: Function UpdateAffectedVertices ( $G_u, T, Dist,$ 
    $Parent$ ,  $Affected\_Del$  and  $Affected$ ):
2:   while  $Affected\_Del$  has true values do
3:     for each vertex  $v \in V$  such that  $Affected\_Del[v] = true$ 
       in parallel do
4:        $Affected\_Del[v] \leftarrow false$ 
5:       for all vertex  $c$ , where  $c$  is child of  $v$  in the SSSP tree  $T$  do
6:         Set  $Dist[c]$  as infinity
7:          $Affected\_Del[c] \leftarrow True$ 
8:          $Affected[c] \leftarrow True$ 
9:     while  $Affected$  has true values do
10:      for each vertex  $v \in V$  such that  $Affected[v] = true$  in parallel do
11:         $Affected[v] \leftarrow False$ 
12:        for vertex  $n$ , where  $n \in V$  and  $n$  is neighbor of  $v$  do
13:          if  $Dist[n] > Dist[v] + W(v, n)$  then
14:             $Dist[n] \leftarrow Dist[v] + W(v, n)$ 
15:             $Parent[n] \leftarrow v$ 
16:             $Affected[n] \leftarrow True$ 
17:          else if  $Dist[v] > Dist[n] + W(n, v)$  then
18:             $Dist[v] \leftarrow Dist[n] + W(n, v)$ 
19:             $Parent[v] \leftarrow n$ 
20:             $Affected[v] \leftarrow True$ 
```

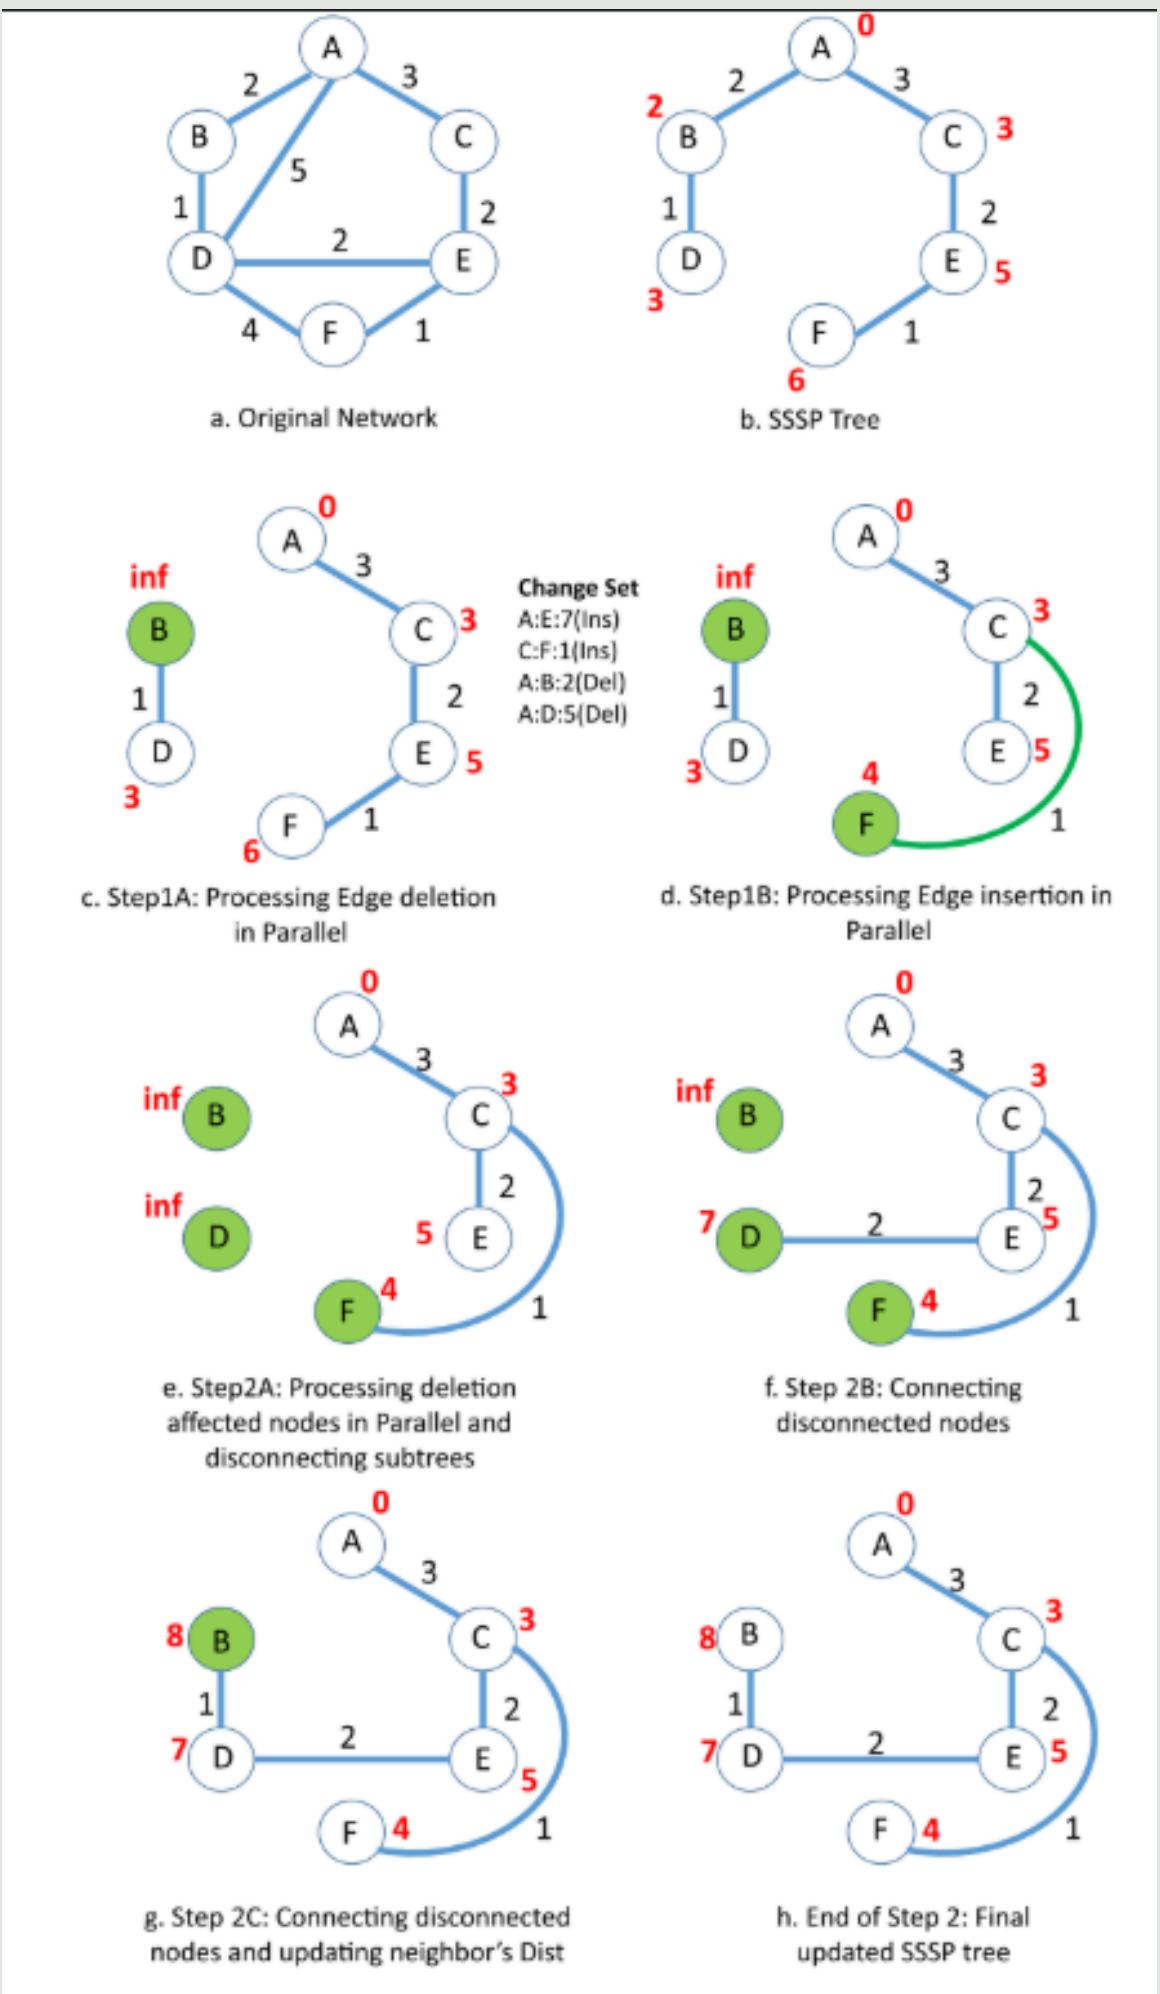


Fig. 2. Illustration of updating SSSP tree.

GPU & CPU IMPLEMENTATION

- OpenMP (CPU): dynamic scheduling, async updates, batches
- CUDA (GPU):
 - Uses VMFB (Vertex-Marking Functional Block)
 - Reduces atomics with a 3-stage marking process
 - Graph stored in CSR format

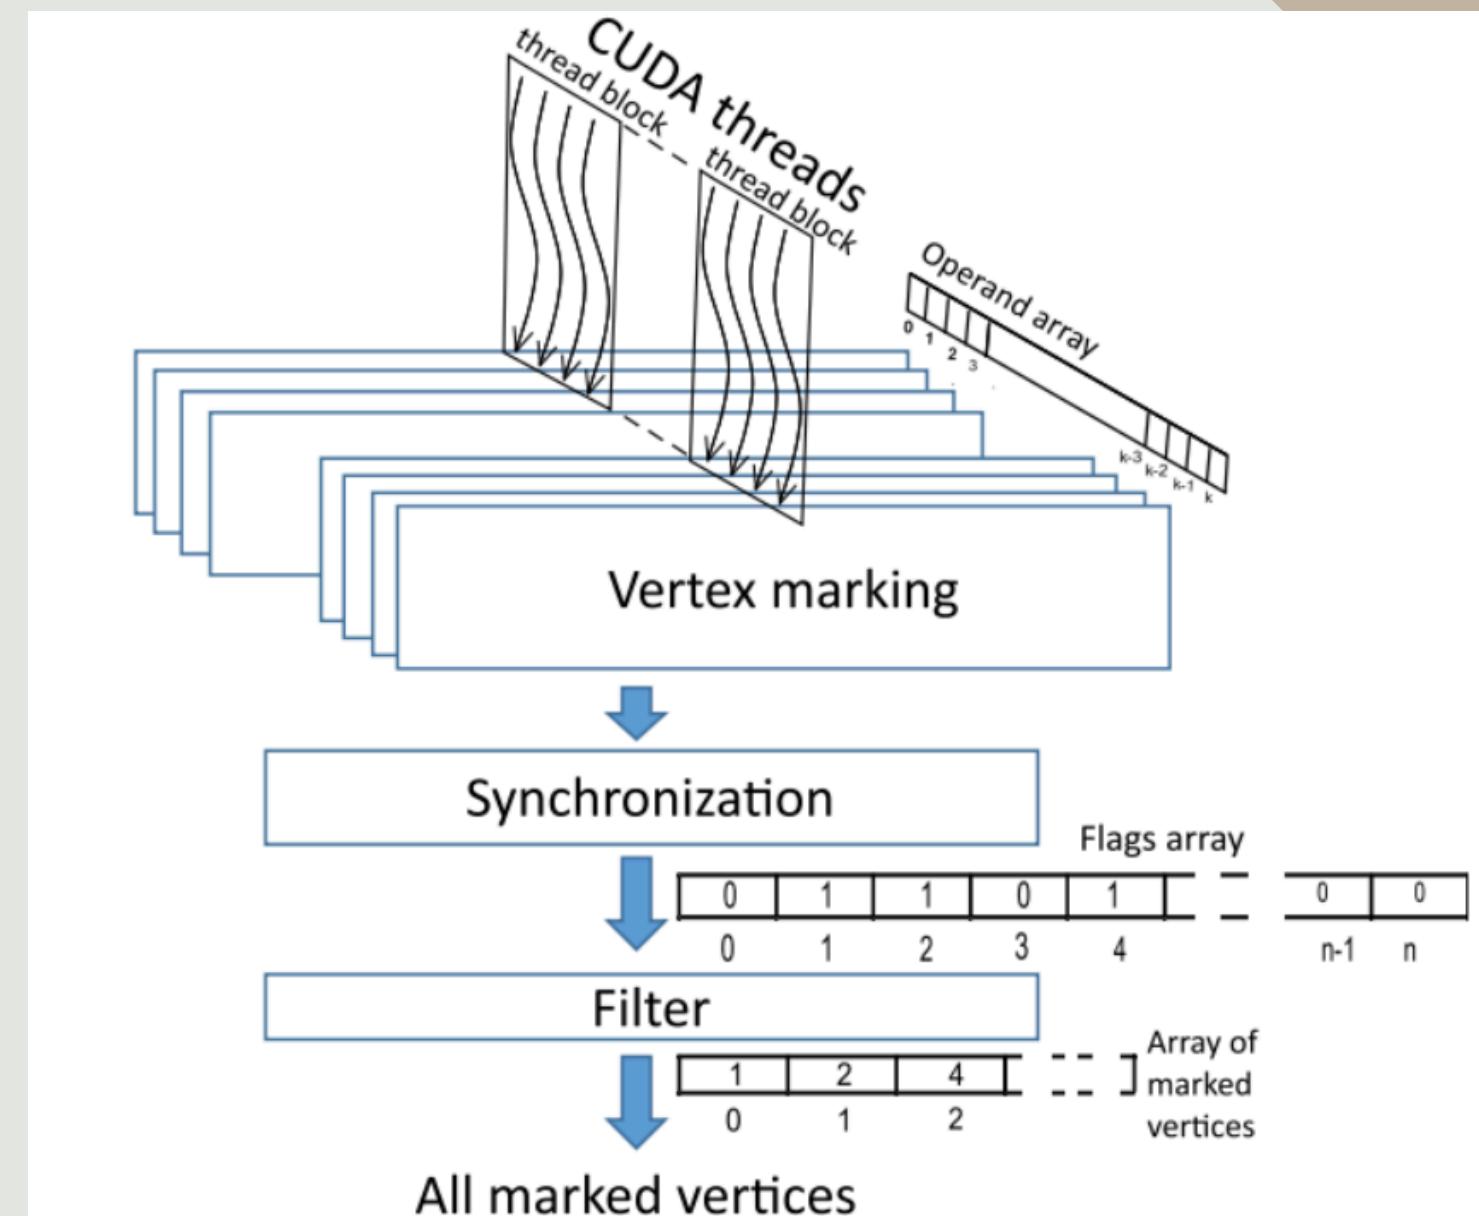


Fig. 4. Vertex-marking functional block (VMFB).

EXPERIMENTAL RESULTS

- Tested on 6 graphs (up to 16M nodes, 258M edges)
- GPU (Tesla V100): up to 8.5x faster than Gunrock
- CPU (Intel Xeon, 72 threads): up to 5x faster than Galois
- Performs best when insertions \geq 50% of updates

TIME & SPACE COMPLEXITY

- Step 1: $O(m/p)$ – m = changes, p = processors
- Step 2: $O(Dxd/p)$ – D = diameter, x = affected vertices, d = avg. degree
- Memory: $O(V)$ for all data structures (Parent, Dist, Affected)

PARALLELIZATION STRATEGY - OVERVIEW

- Enhance using 3 parallel computing tools:
 - **MPI**: inter-node communication
 - **OpenMP/OpenCL**: intra-node processing
 - **METIS**: graph partitioning

MPI - INTER-NODE PARALLELISM

- Use **MPI** to distribute graph partitions across multiple nodes
- Each node works independently on local subgraphs
- Exchange border vertex updates with neighbors
- Ideal for large graphs on HPC clusters
- Reduces global communication by limiting to boundary vertices
- Integrates well with METIS-partitioned graph structure

OPENMP / OPENCL - INTRA-NODE PARALLELISM

- **OpenMP:**
 - Efficient for shared-memory systems
 - Use threads to process affected vertices/subtrees concurrently
 - Dynamic scheduling handles load imbalance
- **OpenCL:**
 - Platform-independent GPU support (NVIDIA, AMD, Intel)
 - Port CUDA kernels for cross-device compatibility
 - Effective for marking, filtering, and distance updates

METIS - GRAPH PARTITIONING

- METIS partitions graph into subgraphs with minimal edge cuts
- Balances vertex count across partitions
- Ensures locality and load balancing in MPI execution
- Helps reduce inter-node communication
- Required preprocessing step for distributed SSSP update

WHEN TO USE WHICH STRATEGY

| <u>Environment</u> | <u>Recommended Strategy</u> |
|---------------------------|------------------------------------|
| Distributed cluster | MPI + METIS |
| Multicore server | OpenMP |
| GPU system | CUDA/OpenCL |
| Mixed updates | Hybrid decision engine |

CONCLUSION & FUTURE WORK

- Efficient parallel SSSP update for dynamic graphs
- Reduces overhead vs. recomputation
- Platform-agnostic: runs on CPU and GPU
- Future Work:
 - Hybrid recompute/update logic
 - Prediction models for updates
 - Extend to other metrics (centrality, PageRank)

Thank You

For your attention