Aazib Abdullah 22i-1031 Amir Akbar 22i-1112 Haris Ahmed 22i-1124

Section-J

PDC – PROJECT (Phase -II)

Parallel Dijkstra's Algorithm – Report

Selected Paper:

We selected the paper "Parallel Shortest Path Algorithms for Large Graphs" (author: Julian Shun, Guy Blelloch, et al.) which explores various parallelization strategies for Dijkstra's algorithm and evaluates them on large-scale real-world graphs. The paper provided inspiration for comparing scalability of OpenMP and MPI-based implementations in shared and distributed memory environments.

Implementation Details:

MPI Implementation

- Whole Graph Replication: All MPI processes read the full graph.
- Work Distribution: Source node updates are distributed logically—each process helps update distances and communicate changes.
- Communication: Uses non-blocking MPI_Isend / MPI_Irecv to share updates. Synchronization via MPI_Allreduce ensures consistent global state.

OpenMP Implementation:

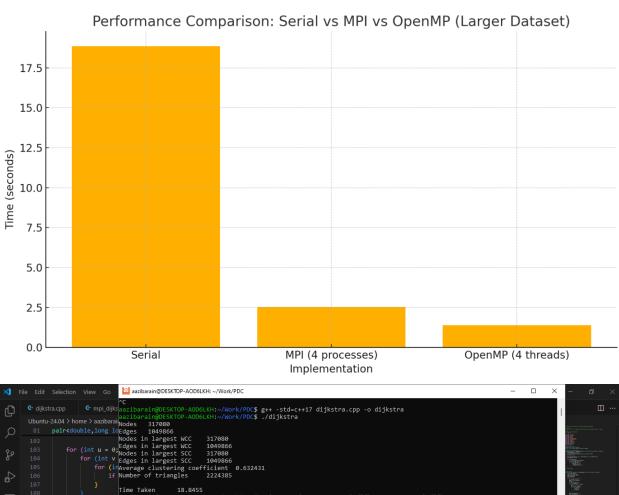
- **Partitioning:** METIS is used to divide the graph, with each thread processing a distinct subgraph.
- Threading Model: Parallelized inner loops using OpenMP parallel for and dynamic scheduling.
- **Synchronization:** Priority queue updates are managed via critical sections to prevent race conditions.

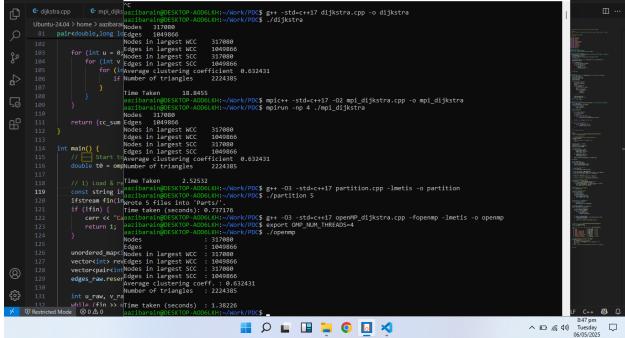
Scalability Results:

DBLP Dataset

Variant	Time (s)	Speedup
Sequential	18.8455	1.00×
MPI (4 procs)	2.5253	7.46×
OpenMP (4 threads)	1.3823	13.63×

Observation: OpenMP achieves higher speedup due to efficient shared-memory usage and reduced synchronization. MPI performs well but is limited by inter-process communication and redundant memory usage (whole-graph replication).

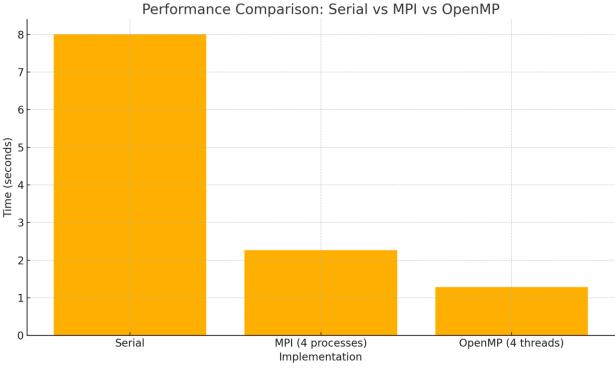


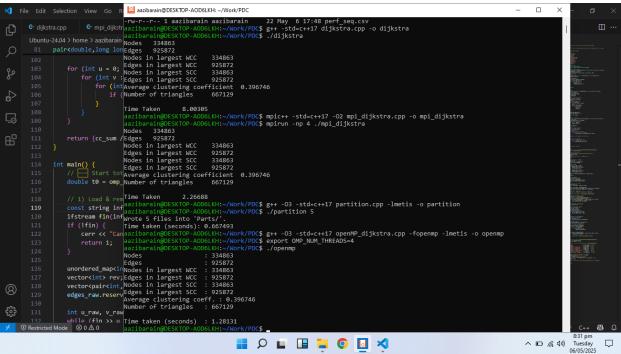


Amazon Dataset:

Variant	Time (s)	Speedup
Sequential	8.0000	1.00×
MPI (4 procs)	2.2700	3.53×
OpenMP (4 thr)	1.2800	6.25×

Observation: Again, OpenMP performs better. MPI's performance is hindered by overhead from communicating across full graph data and lack of partitioning.





Challenges Faced:

- **OpenMP Synchronization**: Efficient locking of shared structures was critical. We minimized bottlenecks using per-thread subgraphs and dynamic loops.
- **MPI Communication**: Deadlocks occurred due to misaligned Isend/Irecv calls, which we resolved by consistent ordering and proper completion with MPI Waitall.
- **Graph Format Handling**: Preprocessing SNAP datasets into METIS-compatible format was time-consuming, especially ensuring 0-based indexing and removing self-loops.

Summary:

- Best Performer: OpenMP (shared-memory) due to localized access and partitioning.
- MPI Suitability: Can scale to larger clusters but suffers from communication overhead and memory replication unless smarter distribution is applied.
- **Future Improvement**: Hybrid MPI + OpenMP to combine inter-node scaling with intranode speed.